

CSCI 420 Computer Graphics

Lecture 7

Shaders

Shading Languages

GLSL

Vertex Array Objects

Vertex Shader

Fragment Shader

[Angel Ch. 1, 2, A]

Oded Stein

University of Southern California

Introduction

- The major advance in real time graphics has been the *programmable* pipeline:
 - First introduced by NVIDIA GeForce 3 (in 2001)
 - Supported by all modern high-end commodity cards
 - NVIDIA, AMD, Intel
 - Software Support
 - Direct3D
 - OpenGL
- This lecture:
programmable pipeline and shaders

OpenGL Extensions

- Initial OpenGL version was 1.0
- Current OpenGL version is 4.6
- As graphics hardware improved, new capabilities were added to OpenGL
 - multitexturing
 - multisampling
 - non-power-of-two textures
 - shaders
 - and many more

OpenGL Grows via Extensions

- Phase 1: vendor-specific: GL_NV_multisample
- Phase 2: multi-vendor:
GL_EXT_multisample
- Phase 3: approved by OpenGL's review board
GL_ARB_multisample
- Phase 4: incorporated into OpenGL (v1.3)

OpenGL 2.0 Added Shaders

- Shaders are customized programs that replace a part of the OpenGL pipeline
- They enable many effects not possible by the fixed OpenGL pipeline
- Motivated by Pixar's Renderman (offline shader)

Shaders Enable Many New Effects



Complex materials



Shadows




Lighting environments



Advanced mapping

Vertex Shader





5x5 terrain (as in hw1)

5x5 = 25 vertices

4x4 = 16 quads

Vertex Shader




96 vertex shaders
execute in parallel

User must tessellate into
triangles (in the VBO)

$4 \times 4 \times 2 = 32$ triangles


$32 \times 3 = 96$ vertices (assuming GL_TRIANGLES)

Fragment Shader



This triangle
rasterizes into
13 pixels

Rasterization

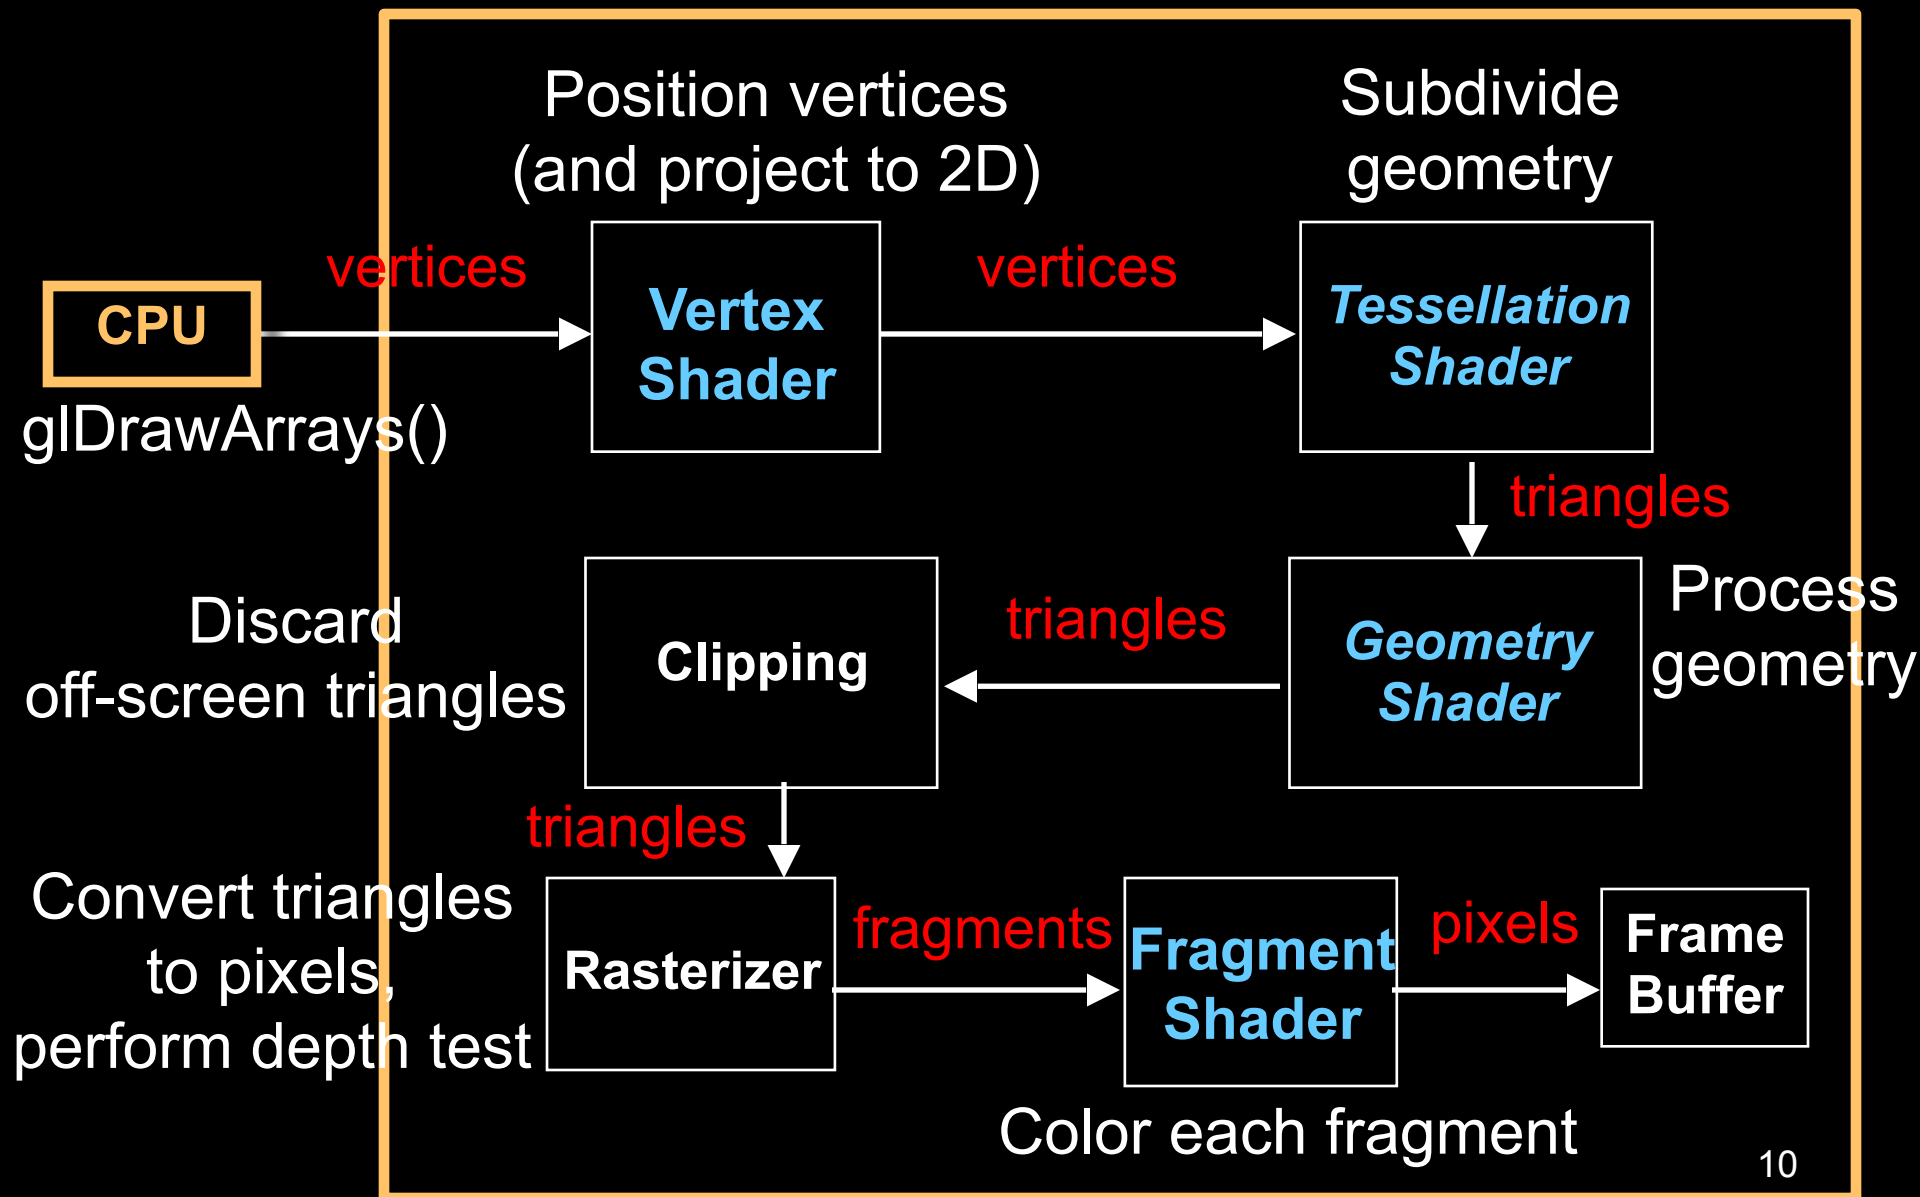


13 fragment shaders
execute in parallel

Some pixels may repeat in multiple triangles

The Rendering Pipeline

GPU



Shaders

- **Vertex shader** (= vertex program)
- **Tessellation control and evaluation shader** (OpenGL 4.0; subdivide the geometry)
- **Geometry shader** (OpenGL 3.2; process, generate, replace or delete geometry)
- **Fragment shader** (= fragment program)
- **Compute shader** (OpenGL 4.3; general purpose)

Shaders

- Compatibility profile: Default shaders are provided by OpenGL (*fixed-function pipeline*)
- Core profile: no default vertex or fragment shader; must be provided by the programmer
- Tessellation shaders, geometry shaders and compute shaders are *optional*

Shader Variables Classification

- Attribute

- Information specific to each vertex/pixel passed to vertex/fragment shader

Example:
Vertex Color

- Uniform

- Constant information passed to vertex/fragment shader
- Cannot be written to in a shader

Example:
Light Position
Eye Position

- Out/in

- Info passed from vertex shader to fragment shader
- Interpolated from vertices to pixels
- Write in vertex shader, but only read in fragment shader

Example:
Vertex Color
Texture Coords

- Const

- To declare non-writable, constant variables

Example:
pi, e, 0.480

Shaders Are Written in *Shading Languages*

- Early shaders: assembly language
- Since ~2004: high-level shading languages
 - OpenGL Shading Language (GLSL)
 - highly integrated with OpenGL
 - Cg (NVIDIA and Microsoft), very similar to GLSL
 - HLSL (Microsoft), the shading language of Direct3D
 - All of these are simplified versions of C/C++

GLSL

- The shading language of OpenGL
- Managed by OpenGL Architecture Review Board
- Introduced in OpenGL 2.0
- We use shader version 1.50:
#version 150
(a good version supporting the core profile features)
- Current shader version: 4.60
 - (it's been that for a while - 2018)

Vertex Shader

- Input: **vertices**, in object coordinates, and per-vertex attributes:
 - color
 - normal
 - texture/UV coordinates
 - many more
- Output:
 - vertex location in clip coordinates
 - vertex color
 - vertex normal
 - many more are possible

Basic Vertex Shader in GLSL

```
#version 150

in vec3 position; // input position, in object coordinates
in vec4 color; // input color
out vec4 col; // output color

uniform mat4 modelViewMatrix; // uniform variable to store the modelview mtx
uniform mat4 projectionMatrix; // uniform variable to store the projection mtx

void main()
{
    // compute the transformed and projected vertex position (into gl_Position)
    gl_Position = projectionMatrix * modelViewMatrix * vec4(position, 1.0f);
    // compute the vertex color (into col)
    col = color;
}
```

Fragment Shader

- Input: **fragments** (tentative pixels), and per-pixel attributes:
 - color
 - normal
 - texture coordinates
 - many more are possible
- Inputs are outputs from the vertex shader, interpolated (by the GPU) to the pixel location !
- Output:
 - pixel color
 - depth value
 - can discard the fragment using the **discard** keyword

Basic Fragment Shader

```
#version 150

in vec4 col; // input color (computed by the interpolator)
out vec4 c; // output color (the final fragment color)

void main()
{
    // compute the final fragment color
    c = col;
}
```

Another Fragment Shader

```
#version 150

in vec4 col; // input color (computed by the interpolator)
out vec4 c; // output color (the final fragment color)

void main()
{
    // compute the final fragment color
    c = vec4(1.0, 0.0, 0.0, 1.0);
}
```

Perspective projection

- Where do we handle the final perspective divide?
- Is it even possible in a shader? Or should OpenGL handle it for us?
- Exercise!

Pipeline program

- Container for all the shaders
- Vertex, fragment, geometry, tessellation, compute
- Can have several pipeline programs
(for example, one for each rendering style)
- Must have at least one (core profile)
- At any moment of time, exactly one pipeline program is bound (active)

Installing Pipeline Programs

- Step 1: Create Shaders
 - Create handles to shaders
- Step 2: Specify Shaders
 - load strings that contain shader source
- Step 3: Compiling Shaders
 - Actually compile source (check for errors)
- Step 4: Creating Program Objects
 - Program object controls the shaders
- Step 5: Attach Shaders to Programs
 - Attach shaders to program objects via handle
- Step 6: Link Shaders to Programs
 - Another step similar to attach
- Step 7: Enable Shaders
 - Finally, let OpenGL and GPU know that shaders are ready

Our helper library: PipelineProgram

```
// load shaders from a file
int BuildShadersFromFiles(const char * filenameBasePath,
    const char * vertexShaderFilename,
    const char * fragmentShaderFilename,
    const char * geometryShaderFilename = NULL,
    const char * tessellationControlShaderFilename = NULL,
    const char * tessellationEvaluationShaderFilename = NULL);
```


Our helper library: PipelineProgram

```
// load shaders from a C text string
int BuildShadersFromStrings(const char * vertexShaderCode,
    const char * fragmentShaderCode,
    const char * geometryShaderCode = NULL,
    const char * tessellationControlShaderCode = NULL,
    const char * tessellationEvaluationShaderCode = NULL);
```

Setting up the Pipeline Program

```
// global variable
PipelineProgram pipelineProgram;

// during initialization:
pipelineProgram.BuildShadersFromFiles("../openGLHelper",
    "vertexShader.glsl", "fragmentShader.glsl");

// before rendering, bind (activate) the pipeline program:
pipelineProgram.Bind();

If you want to use a different pipeline program,
then "Bind" that other pipeline program.
```

Setting up the Uniform Variables

Uploading the modelview matrix transformation to the GPU
(in the display function)

```
float m[16]; // column-major
// here, must fill m (missing code; use OpenGLMatrix class)
// ...

// upload m to the GPU
pipelineProgram.Bind();
GLboolean isRowMajor = GL_FALSE;
pipelineProgram->SetUniformVariableMatrix4fv(
    "modelViewMatrix", isRowMajor, m);
```

Setting up the Uniform Variables



Repeat the same process also for the projection matrix:

```
float p[16]; // column-major
// here, must fill p (missing code; use OpenGLMatrix class)
// ...



// upload p to the GPU
pipelineProgram.Bind();
GLboolean isRowMajor = GL_FALSE;
pipelineProgram->SetUniformVariableMatrix4fv(
    "projectionMatrix", isRowMajor, p);
```

Vertex Array Objects (VAOs)

- A container to collect the VBOs of each object and connect each shader variable with a VBO



VAO 1



VAO 2

shader
variables

Vertex Array Objects (VAOs)

- A container to collect the VBOs of each object
- Usage is mandatory (by the OpenGL standard)
- During initialization:
 - create VBOs (one or more per object),
 - create VAOs (one per object),
 - place the VBOs into the VAO, and connect VBOs to shader variables
- At render time: bind the VAO, then call `glDrawArrays()`

VAO code (initialization)

During initialization:

```
// Create a VAO.  
VAO * vao = new VAO();  
// Connect the shader variables to their respective VBOs.  
vao->ConnectPipelineProgramAndVBOAndShaderVariable(  
    pipelineProgram, vboPositions, "position");  
vao->ConnectPipelineProgramAndVBOAndShaderVariable(  
    pipelineProgram, vboColors, "color");
```

Using the VAO for rendering

In the display function:

```
// Bind the vertex and fragment shaders to use.
pipelineProgram->Bind();
// Select which object to render.
vao->Bind();

// Render the object contained in the VAO.
GLint first = 0;
GLsizei count = numVertices;
glDrawArrays(GL_TRIANGLES, first, count);
```


GLSL: Standard

- GLSL is a well-specified standard
 - <https://www.khronos.org/opengles/sdk/docs/manglsl/docbook4/>
 - Really useful online reference
- Doesn't always work as specified... depends on the driver vendor.
- Don't write crazy C code using syntax features.
 - But do use crazy math!

GLSL: Data Types

- Scalar Types
 - float - 32 bit, very nearly IEEE-754 compatible
 - int - at least 16 bit
 - bool - like in C++
- Vector Types
 - vec[2 | 3 | 4] - floating-point vector
 - ivec[2 | 3 | 4] - integer vector
 - bvec[2 | 3 | 4] - boolean vector
- Matrix Types
 - mat[2 | 3 | 4] - for 2x2, 3x3, and 4x4 floating-point matrices
- Sampler Types
 - sampler[1 | 2 | 3]D - to access texture images

GLSL: Data Types

- Structs
 - Very useful to shuffle specific information between CPU and GPU in one block.
 - Careful - padding!

```
struct Light {  
    float intensity;  
    vec3 position;  
    vec4 color;  
};
```

```
.....
```

```
uniform Light l1;
```

```
.....
```

```
void main()  
{  
    Light l = ...;  
}
```

GLSL: Operations

- Operators behave like in C++
- Component-wise for vector & matrix
- Multiplication on vectors and matrices
- Examples:
 - `Vec3 t = u * v;`
 - `float f = v[2];`
 - `v.x = u.x + f;`

GLSL: Swizzling

- Swizzling is a convenient way to access individual vector components

```
vec4 myVector;  
myVector.rgba; // is the same as myVector  
myVector.xy; // is a vec2  
myVector.b; // is a float  
myVector[2]; // is the same as myVector.b  
myVector.xb; // illegal  
myVector.xxx; // is a vec3
```

0123 ~ xyzw

GLSL: Flow Control

- Loops
 - C++ style if-else
 - C++ style for, while, and do

Example:

```
if(isHidden) {  
    color.w = 0.;  
}
```

...

```
for(int i=0; i<numLights; ++i)  
{  
    color += light[i].c;  
}
```

GLSL: Flow Control

- Loops
 - C++ style if-else
 - C++ style for, while, and do
- Jumps
 - `continue;`
 - `break;`
 - `return;`
 - `discard;` (in fragment shader)

GLSL: Flow Control

- Functions
 - Much like C++
 - Entry point into a shader is void main()
 - No support for recursion
 - Call by value-return calling convention

Example:

```
float sign(float x) {  
    if(x>0) {  
        return 1.;  
    } else if(x<0) {  
        return -1.;  
    } else {  
        return 0.;  
    }  
}
```


GLSL: Flow Control

- Functions
 - C++-like overloading exists.
 - Arguments can be `const`
 - You can add a parameter qualifier
 - `in` (default; will be copied into variable on function call)
 - `out` (will be copied into variable on function return)
 - `inout` (will be copied into variable on function call and return)
 - `const` contradicts `out` and `inout`
 - Poor man's pass-by-reference
 - Not actual references, only copying.
 - Avoid if you can because of confusion.

GLSL: Flow Control

Example:

```
void sign(in float x, out float y) {  
    if(x>0) {  
        y = 1.;  
    } else if(x<0) {  
        y = -1.;  
    } else {  
        y = 0.;  
    }  
}
```

GLSL: Built-in Functions

- Wide Assortment
 - Trigonometry (cos, sin, tan, etc.)
 - Exponential (pow, log, sqrt, etc.)
 - Common (abs, floor, min, clamp, etc.)
 - Geometry (length, dot, normalize, reflect, etc.)
 - Relational (lessThan, equal, etc.)
- Need to watch out for common reserved keywords
- Always use built-in functions, do not implement your own
 - Sign function from last slide is a bad idea.
- Some functions are not implemented on some cards

GLSL: Built-in Variables

- Always prefaced with `gl_`
- Accessible to both vertex and fragment shaders
- Examples:
 - (input) `gl_VertexID`: index of currently processed vertex
 - (input) `gl_FrontFacing`: whether pixel is front facing or not
 - (input) `gl_FragCoord` : x,y: coordinate of pixel, z: depth
 - (output) `gl_FragDepth`: pixel depth

Debugging Shaders

- More difficult than debugging C programs
- Common show-stoppers:
 - Typos in shader source
 - Assuming implicit type conversion (cannot convert vec4 to vec3)
 - Attempting to connect VAOs to non-existent (say, due to a typo) shader variables
 - Driver bugs beyond your control
 - <https://stackoverflow.com/a/23999304>
- Very important to check error codes; use status functions like:
 - `glGetShaderiv(GLuint shader, GLenum pname, GLint * params)`

A few debugging best practices

- printf debugging corresponds to just writing into an output color

```
if(condition) {  
    col = vec4(1., 0., 0., 1.);  
} else {  
    col = vec4(0., 1., 0., 1.);  
}
```

- Pass variables through to fragment shader.
- Display input variables.
- Write multiple functions and color-debug them separately.

Summary

- Shading Languages
- Program Pipeline
- Vertex Array Objects
- GLSL
- Vertex Shader
- Fragment Shader

Shadertoy!

- www.shadertoy.com
- Website that allows you to write a fragment shader and run it in the web
- Fragment shaders are very powerful. You can do a lot with them

Demoscene

- Computer graphics subculture
- Writing small programs that show impressive visual videos
- File size restriction (1MB, 64KB, ...), old hardware...
- You can do a lot with just generating images with math, you don't need to save/load assets!
- Beautiful art form

Demoscene

- Running on 1MB memory Amiga
- Lotus <https://www.pouet.net/prod.php?which=81094>



Shadertoy!

- Basic example:
 - <https://www.shadertoy.com/view/MlySzw>
- More complicated examples:
 - <https://www.shadertoy.com/view/MsjXDm>
 - <https://www.shadertoy.com/view/mtyGWy>
- Crazy cool examples (too complicated for me to understand...):
 - <https://www.shadertoy.com/view/4td3zj>
 - <https://www.shadertoy.com/view/MfjyWK>
 - <https://www.shadertoy.com/view/lsXGzH>
 - <https://www.shadertoy.com/view/lsf3zr>